

CVPR 2020 Tutorial on

From HPO to NAS: Automated Deep Learning

Hang Zhang^[1], Cedric Archambeau^[1], Song Han^[2], Matthias Seeger^[1],
Thibaut Lienart^[1], Chongruo Wu^[1,3], Ligeng Zhu^[2], Haotian Tang^[2], Aaron Klein^[1]

Amazon¹, MIT², UC Davis³



Automated Deep Learning Tutorial Overview

- AutoML Background and Overview (by Cedric Archambeau)
- Introducing AutoGluon Toolkit (by Hang Zhang)
- AutoML for TinyML with Once-for-All Network (by Song Han)
- Hands-on Section
 - Quick Start and ProxylessNAS Tutorial (by Chongruo Wu)
 - Advanced HPO Algorithm (by Thibaut Lienart)
 - Once-For-All Network Tutorial (by Ligeng Zhu and Haotian Tang)
 - Code Example Walk-through (by Hang Zhang)

<https://hangzhang.org/CVPR2020/>
<https://autogluon.mxnet.io/>

Auto**G**LUON: An AutoML Toolkit for Deep Learning

Hang Zhang

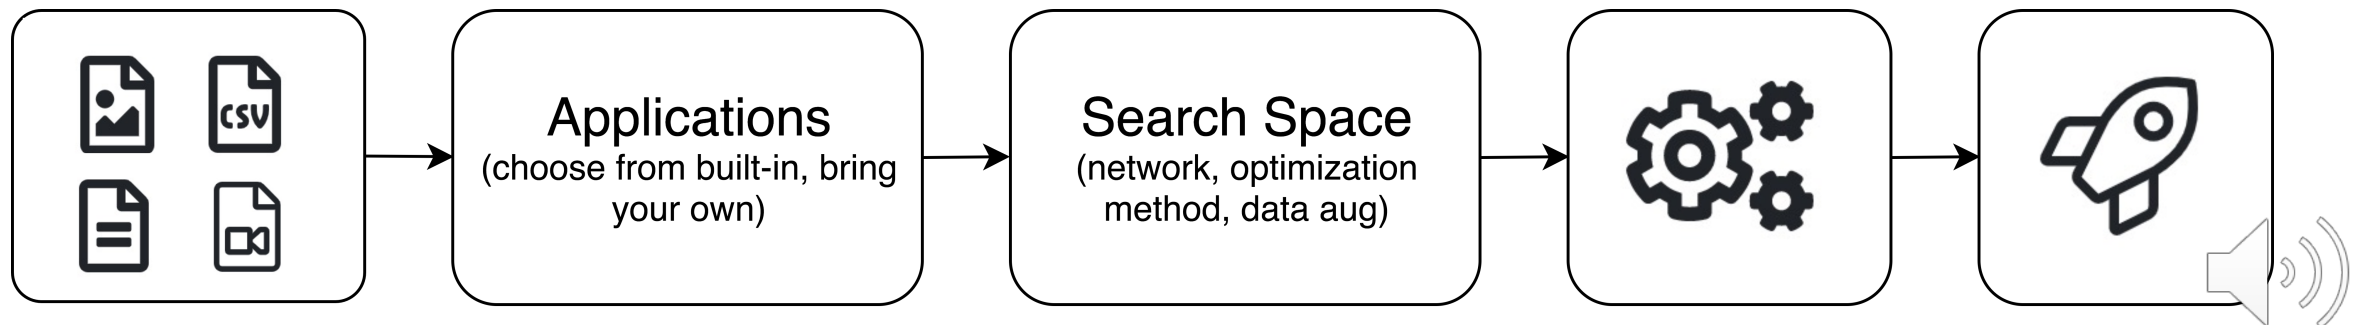
Applied Scientist, Amazon

CVPR Tutorial on From HPO to NAS: Automated Deep Learning (Jun 2020)




Auto LUON

- A Popular Toolkit:
 - 2.3K GitHub Starts, many media posts online
- AutoGluon Users:
 - DL expert/researcher (improves their model performance)
 - Researcher on AutoML/HPO algorithm
 - Developer conducting large scale distributed training (GPU management)
- Research Prototype => AutoML Production Pipeline (SageMaker)

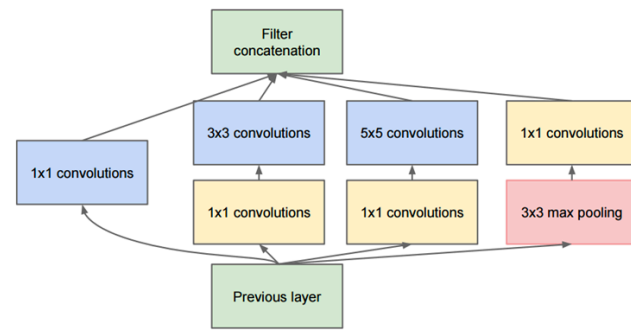


Why HPO in Computer Vision?


- Feature Engineering => Network Engineering => Architecture Space Design
- Increasing number of design choices



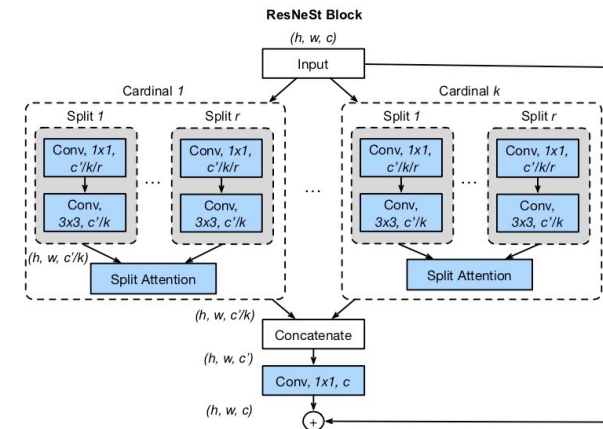
Very Deep Convolutional Networks for Large-Scale Image Recognition




Going Deeper with Convolutions



Squeeze-and-Excitation Networks



ResNeSt: Split-Attention Network

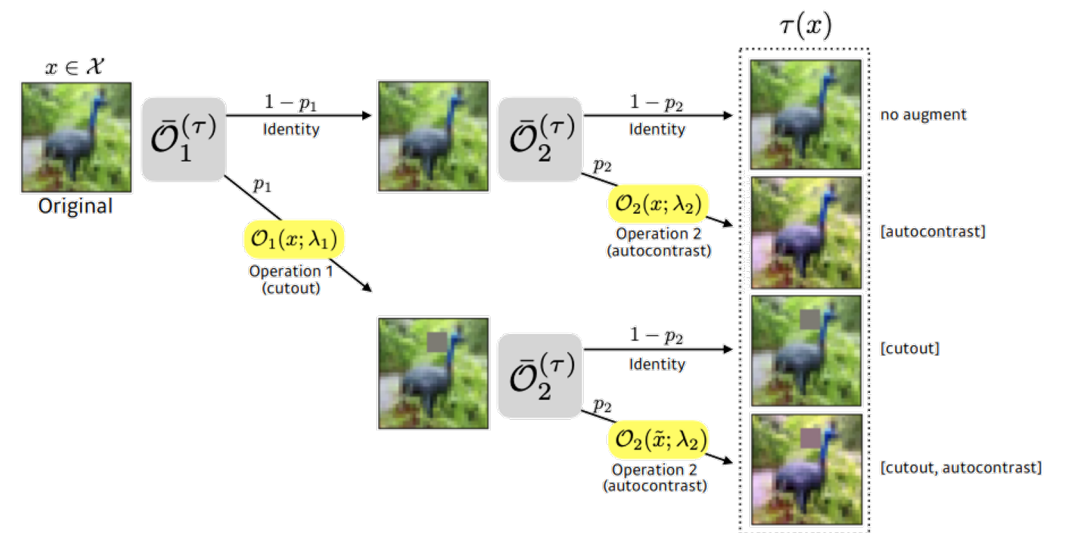


ProxlessNAS: Direct Neural Architecture Search on Target Task and Hardware



HPO for Computer Vision

- Optimization of training strategies and HPs
 - AutoAugment (Fast AA, RandAA)
 - FBNetV3 (joint optimization)
- Train recipe:
 - Data augmentation policies
 - Mix-up training
 - Label smoothing
 - Optimization HPs: learning rate, weight decay



What Does HPO Look Like

- Spinning up the experiments using different choices
- Learn lessons and suggest next choice

```
# change the rank for worker node
python train_dist.py --dataset imagenet --model resnet50 --lr-scheduler cos --epochs 270 --
checkname resnet50 --lr 0.025 --batch-size 64 --dist-url tcp://MASTER:NODE:IP:ADDRESS:23456 --
world-size 4 --label-smoothing 0.1 --mixup 0.2 --no-bn-wd --last-gamma --warmup-epochs 5 --rand-aug
--rank 0
```



```
# change the rank for worker node
P) # change the rank for worker node
cl # change the rank for worker node
w cl python train_dist.py --dataset imagenet --model resnet50 --lr-scheduler cos --epochs 270 --
- w checkname resnet50 --lr 0.025 --batch-size 64 --dist-url tcp://MASTER:NODE:IP:ADDRESS:23456 --
- world-size 4 --label-smoothing 0.1 --mixup 0.2 --no-bn-wd --last-gamma --warmup-epochs 5 --rand-aug
--rank 0
```



What Auto LUON Provides

- AutoGluon Toolkit
 - Easy-to-use and easy-to-customize
 - Efficient in Search (BayesOpt, Hyperband, Reinforcement Learning)
 - Scalable (Distributed Search, AWS Batch, SageMaker Integration)
- Goal of this tutorial
 - Change the conventional way of doing research manually
 - Constructing search space and pass the workload to the machines

Easy to Use – Search Space

- Integer Space (similarly for Categorical and Real Space):

```
import autogluon as ag
a = ag.space.Int(lower=0, upper=10)
print(a)
print(a.rand())
```

```
Int: lower=0, upper=10
```

```
7
```

- Dict Space (similarly for List Space):

```
g = ag.space.Dict(key1=ag.space.Categorical('alpha', 'beta'),
                  key2=ag.space.Int(0, 3),
                  key3='constant')
print(g.rand())
```

```
{'key1': 'beta', 'key2': 1, 'key3': 'constant'}
```

Easy to Customize – Training Function

- Allow HPO on python training scripts
automatically generate search algorithm and conduct HPO

```
@ag. args(  
    batch_size=128,  
    lr=ag.Real(1e-4, 1e-2),  
)  
def train_fn(args, reporter):  
    print(args.lr, args.batch_size)  
    reporter(epoch=0, accuracy=0.9)
```

```
myscheduler = ag.scheduler.FIFOScheduler(train_fn, num_trials=25)  
myscheduler.run()
```

Easy to Customize – Searchable Objects

- Defining search space and constructing dynamic object in one go:

- Searchable object from a function:

```
import gluoncv as gcv
@ag.func(multiplier=ag.Categorical(0.25, 1.0))
def get_mobilenet(multiplier):
    return gcv.model_zoo.MobileNetV2(multiplier=multiplier, classes=4)
```

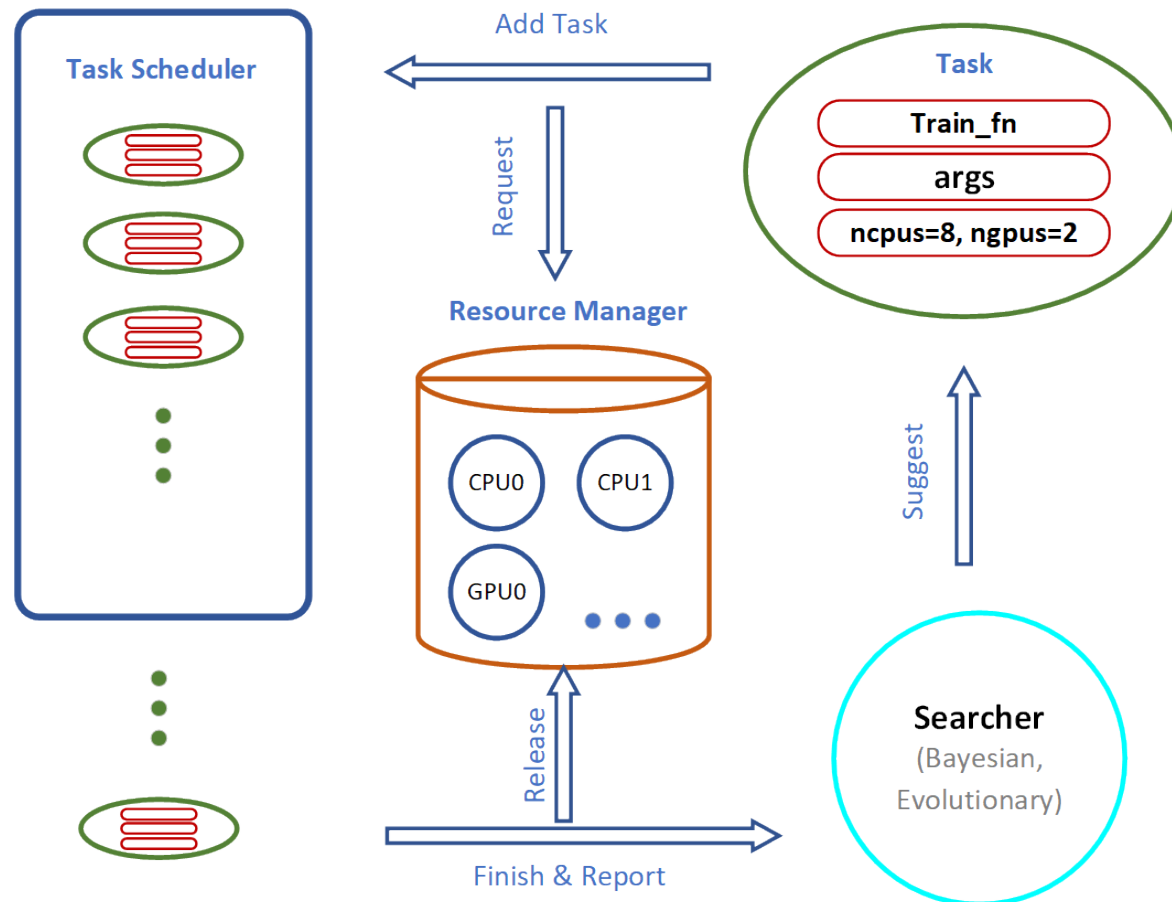
- Searchable object from a class:

```
from mxnet import optimizer as optim
@ag.obj(learning_rate=ag.space.Real(1e-4, 1e-1, log=True),
        wd=ag.space.Real(1e-4, 1e-1))
class Adam(optim.Adam):
    pass
```

- Pass the searchable object as the input of the train_fn:

```
@ag.args(net = mobilenet(), optim = Adam())
def train_fn(args, reporter):
    pass
```

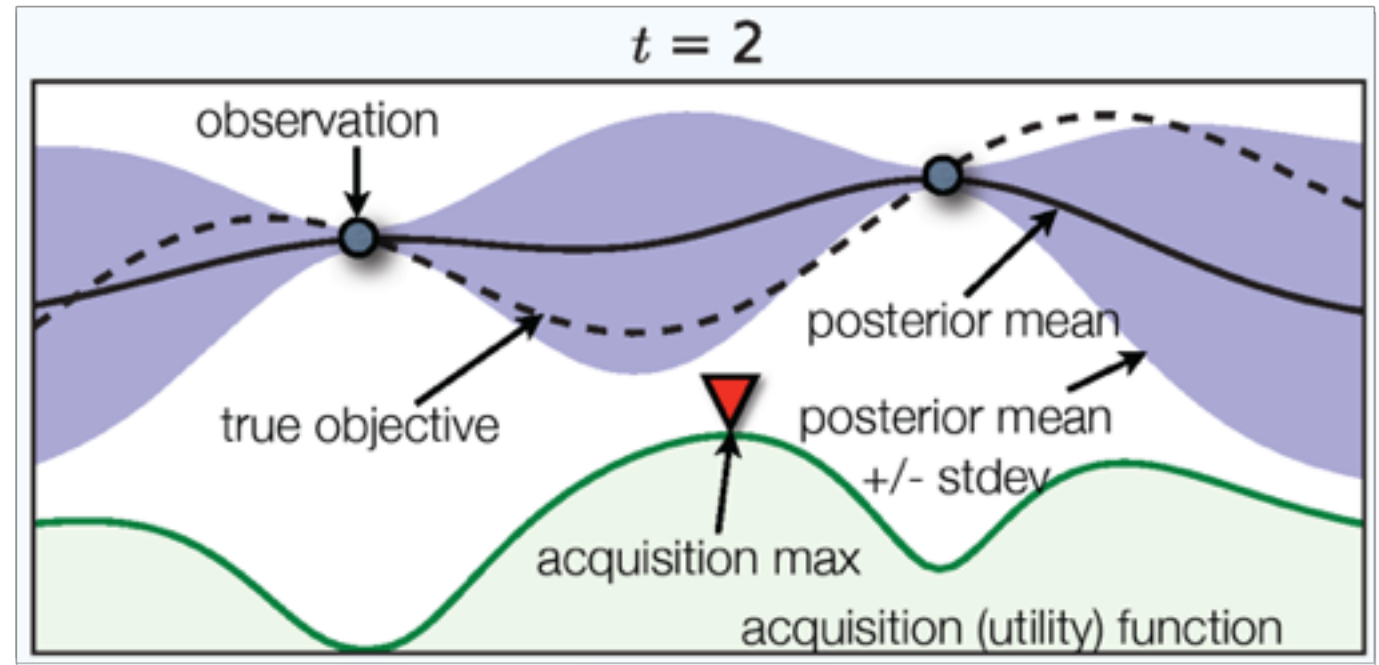
Efficient in Search – Advanced Search Strategies



- Searcher: suggests configurations for the next training jobs.
- Scheduler: schedules the training job when the computation resources are available.
- All components are modularized and easy to extend

Search Algorithms (Bayesian Optimization)

- Surrogate Model:
reflects the probabilistic beliefs
- Acquisition function:
which hyperparameter configuration to try next



```
scheduler = ag.scheduler.FIFOScheduler(train_fn, searcher='bayesopt')
```

<http://autogluon.s3.amazonaws.com/api/autogluon.searcher.html#skoptsearcher>



Search Algorithms (RL, Gluon)

- Auto Construct a LSTM
Controller based on Search Space


$$J(\theta) = E_{a \sim P(a; \theta)}[R(a)]$$

$$\nabla_{\theta} J(\theta) = E_{a \sim P(a; \theta)}[R(a) \nabla_{\theta} \log P(a; \theta)]$$

$$\frac{1}{m} \sum_{k=1}^m \sum_{t=1}^T \log P(a_t | a_{t-1}; \theta) (R_k - b)$$

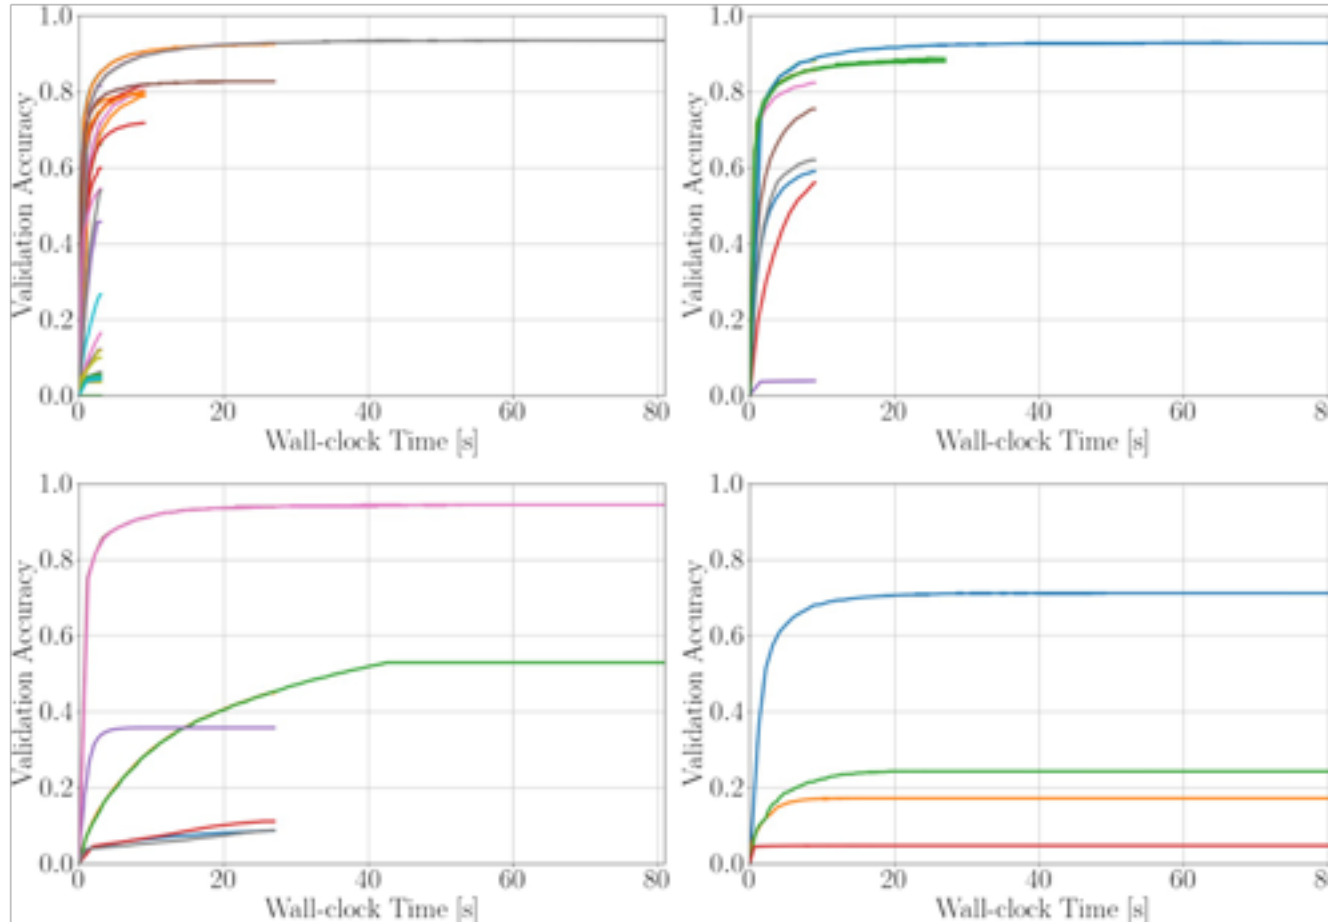


Early Stopping – Successive Halving



- Randomly sample a set of configurations
- Evaluate the performances of all currently remaining configurations
- Throw out, the bottom half of the worst configurations
- Go back to 2. and repeat until one configuration remains.

Early Stopping – Hyperband



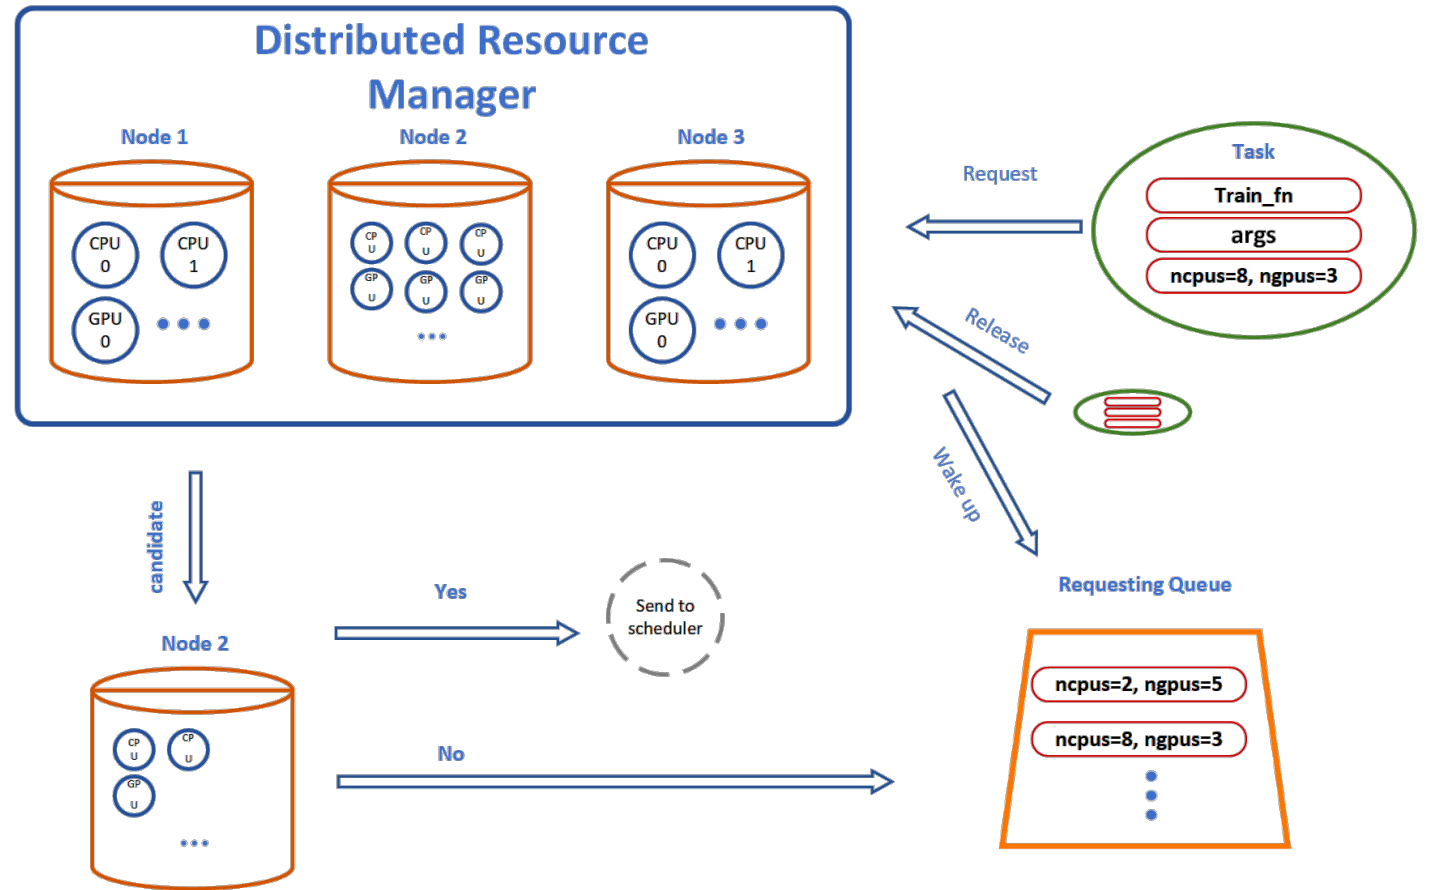
- Number of Total Iters: N
- Number of Successive Halving Brackets: $N_s = \log\left(\frac{N}{\eta}\right) + 1$, where η is a constant
- Number of iters per successive halving: $B = N_s \cdot N$

scheduler = ag.scheduler.
HyperbandScheduler(train_fn)

Scalable -- Distributed Search

- Seamless Experience on Multi-machines
- Automatically manage resource and schedule on remote machines:

```
ip_addrs = ['35.166.23.232',  
            '35.166.30.159']  
results = ag.schedule. FIFOScheduler(  
    train_fn,  
    dist_ip_addrs=ip_addrs)
```



<http://autogluon.s3.amazonaws.com/tutorials/course/distributed.html>

Running on AWS SageMaker

- SageMaker setup all the environment for you
- Upload dataset to S3

```
prefix = 'DEMO-autogluon-cifar10'  
data_location = sess.upload_data('./data', key_prefix=prefix)
```

- Bring your own container:

```
from sagemaker.estimator import Estimator  
estimator = Estimator(image_name=ecr_image,  
                        role=get_execution_role(),  
                        train_instance_count=4,  
                        train_instance_type='p3.2xlarge')  
estimator.fit(data_location)
```

- Download the models from S3

<https://github.com/zhanghang1989/AutoGluon-Docker>

Q&A

Thank you!